

CSE428: Image Processing

Lecture 5

Neighbourhood processing: Part 1

Contents

- Spatial operations
- Neighborhood operations
- Spatial filtering
- Image padding
- Linear spatial filtering
- Correlation and convolution
- Smoothing spatial filters: Gaussian, Average & Median filtering
- Unsharp masking & high boost filtering
- Geometric transformations

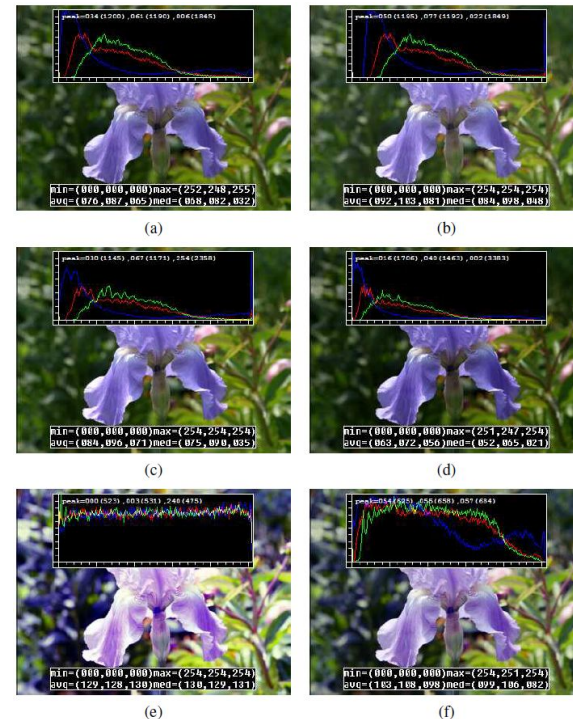
Spatial Operations

Can be subdivided into 3 broad categories:

1. Single-pixel operations or **point processing** (covered in the last lecture)
2. Neighborhood operations or **neighborhood processing**
3. Geometric spatial **transformations** (won't go into too much detail)


Single-pixel operations

- Simplest possible operation on a digital image: alter the pixel intensities
- If T is a transfer function, $s = T(z)$ or
 - $z \rightarrow$ original image pixel intensity $s \rightarrow$ mapped intensity
- Some local image processing operations:
 - (a) original image along with its three color histograms;
 - (b) brightness increased (additive offset, $b = 16$);
 - (c) contrast increased (multiplicative gain, $a = 1.1$);
 - (d) gamma (partially) linearized ($\gamma = 1.2$);
 - (e) full histogram equalization;
 - (f) partial histogram equalization



Neighborhood

- A pixel p at coordinates has four horizontal and vertical neighbors whose coordinates are given by
 - $(x + 1, y)$, $(x - 1, y)$, $(x, y + 1)$, $(x, y - 1)$
 - This set of pixels, called the 4-neighbors of p , $N_4(p)$
- The four diagonal neighbors of p
 - $(x + 1, y + 1)$, $(x + 1, y - 1)$, $(x - 1, y + 1)$, $(x - 1, y - 1)$
 - Diagonal neighbors of p , $N_D(p)$
- 4-neighbors together with the diagonal neighbors are called the 8-neighbors of p : $N_8(p)$



Distance measures

- For pixels p , q , and z , with coordinates (x, y) , (s, t) , and (v, w)
- Euclidean distance between p and q
 - $D_e(p, q) = [(x - s)^2 + (y - t)^2]^{1/2}$
- D_4 distance (city-block distance) between p and q
 - $D_4(p, q) = |x - s| + |y - t|$
 - The pixels with $D_4 = 1$ are the 4-neighbors of (x, y)
- D_8 distance (chessboard distance) between p and q
 - $D_8(p, q) = \max(|x - s|, |y - t|)$
 - The pixels with $D_8 = 1$ are the 8-neighbors of (x, y)

			2		
		2	1	2	
2	1	0	1	2	
		2	1	2	
			2		

2	2	2	2	2
2	1	1	1	2
2	1	0	1	2
2	1	1	1	2
2	2	2	2	2

Linear vs Nonlinear Operations

Consider a general operator, H such that : $H [f(x, y)] = g(x, y)$

H is said to be a linear operator if the following is true:


$$H [a_1 f_1(x, y) + a_2 f_2(x, y)] = a_1 H [f_1(x, y)] + a_2 H [f_2(x, y)]$$

1. *Additivity* property: output of a linear operation due to the sum of two inputs is the same as performing the operation on the inputs individually and then summing the results
2. *Homogeneity* property: the output of a linear operation to a constant times an input is the same as the output of the operation due to the original input multiplied by that constant

Neighborhood operations

S_{xy} is the set of coordinates of a neighborhood centered on an arbitrary point (x, y) in an image f

- *Neighborhood processing* generates a corresponding pixel at the same coordinates in an output (processed) image, g , such that the value of that pixel is determined by a specified operation involving the pixels in the input image with coordinates in S_{xy}



Digital Image Processing, Third Edition,
Rafael C. Gonzalez & Richard E. Woods

Neighborhood operations: Mechanism

Spatial filtering consists of

1. A **neighborhood** (defined by the filter kernel, or mask)
2. A **predefined operation** (which is performed on the neighborhood)
 - Linear operation
 - Nonlinear operation

Filtering creates a new pixel at the same position of the center of the neighborhood, and whose value is the result of the filtering operation

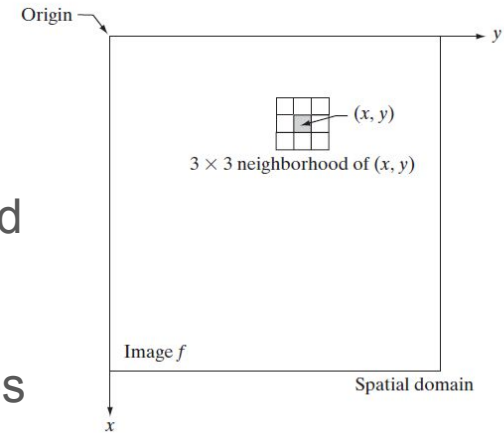
A processed (filtered) image is generated as the center of the filter visits each pixel in the input image

Spatial Filtering

Neighborhood processing is also called ***spatial filtering***

- $g(x, y) = T[f(x, y)]$
- Input image: $f(x, y)$
- Output image: $g(x, y)$
- $T[.]$ is an operator in f , defined over a neighborhood
 - Linear operation -> linear spatial filtering
 - Nonlinear operation -> nonlinear spatial filtering
- T can be applied to a single image or multiple images

Example: a 3x3 neighborhood of (x, y)




Linear Spatial Filtering

Linear spatial filtering in a $m \times n$ neighborhood


- Mask/filter/kernel size: $m \times n$
 - Assume $m = 2a + 1$, $n = 2b + 1$ (both odd numbers)
- Image size: $M \times N$

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$

Example: $g(x, y) = w(-1, -1) f(x-1, y-1) + w(-1, 0) f(x-1, y) + \dots + w(0, 0) f(x, y) + \dots + w(1, 1) f(x+1, y+1)$



But.. what about the border pixels?



Solution 1: Don't visit the border pixels!
Result: Image is ***shrunked***

Solution 2: *Padding*

Result: Image shape *stays the same*

Image Padding

1. **Zero** padding: extend the borders of the original image with zeros
2. **Constant** padding: extend the borders of the original image with a constant pixel value
3. **Mirror** padding: extend the borders of the original image by reflecting the edge pixels
4. **Clamp/Edge/Nearest** padding: extend the borders of the original image by repeating edge pixels

More ways: <https://scikit-image.org/docs/dev/api/skimage.util.html#skimage.util.pad>





Image Padding

A 1448x1448x3 image padded with 200 pixels on each side:

1. Original image
2. Zero padding
3. Constant padding
4. Mirror padding
5. Edge padding
6. Linear(ramp) padding

1. Original image (no padding)




2. Zero padding (constant = 0)




3. Constant padding (constant = 200)



4. Mirror padding



5. Edge padding




6. Linear ramp padding



Linear Spatial Filtering: Demo

- Kernel shape: 3x3
- Input image shape: 6x6
- Output image shape: 6x6
- Padding: Nearest
- Step (stride): 1

- Kernel, $w(x, y) = \begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$




Linear Spatial Filtering: Output Shape

- Kernel shape: $m \times n$
- Input image shape: $M \times N$
- Number of pixels padded: p
- Stride: s
- Output image shape: $H \times W$

$$H = \text{floor}\left(\frac{M + 2p - m}{s} + 1\right)$$

$$W = \text{floor}\left(\frac{N + 2p - n}{s} + 1\right)$$



Correlation

Correlation between two 2D signals $w(x, y)$ & $f(x, y)$ is defined as:

$$w(x, y) \star f(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x+s, y+t)$$

- Outputs a ***measure of similarity*** of the two signals as a function of the displacement of one relative to the other
- Also called the cross-correlation
- Basically a sliding dot product (vector element wise multiplication)
- Used for searching a long signal $[f(x, y)]$ for a shorter signal $[w(x, y)]$ (a known feature)

Linear spatial filtering \Leftrightarrow 2D signal correlation

Correlation to Convolution

- Linear spatial filtering is basically a 2D signal **correlation** operation
- If we rotate the original filter $w(x, y)$ 180° and then proceed, it's called the **convolution** operation
- Mathematically,

$$\text{Correlation: } w(x, y) \star f(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x+s, y+t)$$

$$\text{Convolution: } w(x, y) \star f(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x-s, y-t)$$

For symmetric filters,
convolution and correlation both yield the
same result!

Smoothing Spatial Filters

- Linear
 - Box/average filtering
 - Gaussian filtering
- Non-linear
 - Median filtering

Example: A 3 x 3 filter with coefficients w_1 through w_9

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

Linear Smoothing Filters

1. Averaging filter

- Average of all the pixels encompassed by the filter kernel
- An $m \times n$ averaging *box kernel*: $w(x, y) = 1 / m \times n$

2. Gaussian filter

- Weighted* average of all the pixels encompassed by the gaussian filter kernel
- An $m \times n$ *gaussian kernel*: $w(x, y) = \exp(-(x^2+y^2)/2\sigma^2)$

Example: a 3x3 Box filter and a 3x3 Gaussian filter

- (note that the filters are normalized so that the intensity of the output pixel remains unchanged)

$$\frac{1}{9} \times$$

1	1	1
1	1	1
1	1	1

$$\frac{1}{16} \times$$

1	2	1
2	4	2
1	2	1

Linear Smoothing Filters

Applications

- Image **Blurring** (As the name suggests)
- Image **Sharpening** (wait, sharpening from blurring? what?)
- Image **Denoising**
- **Low pass filtering** (getting rid of high variations)


Gaussian Filtering: The Gaussian Blur

Result of filtering an image with with (25×25) , $\sigma^2 = 8$ Gaussian kernel

Original image




Gaussian blur



“gaussian blur” memes



Kyle
@KylePlantEmoji

Ah  I spilled my gaussian blur



11:45 AM · Oct 16, 2021 · Twitter Web App


Gaussian Filtering: The Gaussian Blur

Increasing the kernel shape increases blurriness, borders become more prominent


(5 x 5)



(55 x 55)



(205 x 205)




Gaussian Filtering vs Average Filtering


For the same kernel size the gaussian blur *preserves more detail* of the image

Even though the image is blurred the edges are somewhat better preserved for the gaussian blurring (edge preserving blurring)


(5 x 5) Box kernel



(55 x 55) Box kernel




(205 x 205) Box kernel




(5 x 5) Gaussian kernel



(55 x 55) Gaussian kernel



(205 x 205) Gaussian kernel



Order Statistic Filtering

Response is based on *ordering* (ranking) the pixels contained in the pixel neighborhood (S_{xy})

- **Median** filtering:

$$\hat{f}(x, y) = \operatorname{median}_{(s,t) \in S_{xy}} \{g(s, t)\}$$

- **Max** and **min** filtering:

$$\hat{f}(x, y) = \max_{(s,t) \in S_{xy}} \{g(s, t)\} \quad \hat{f}(x, y) = \min_{(s,t) \in S_{xy}} \{g(s, t)\}$$

- **Midpoint** filtering:

$$\hat{f}(x, y) = \frac{1}{2} \left[\max_{(s,t) \in S_{xy}} \{g(s, t)\} + \min_{(s,t) \in S_{xy}} \{g(s, t)\} \right]$$

- **Alpha-trimmed** mean filtering: $\hat{f}(x, y) = \frac{1}{mn - d} \sum_{(s,t) \in S_{xy}} g_r(s, t)$

Median Filtering

Median Filter (Order-statistic filter)

- nonlinear spatial filter
- each pixel value is replaced by the median of the intensity values in the neighborhood
- excellent noise-reduction capabilities
- less blurring than linear filters
- particularly effective with impulse noise (salt and pepper noise)

Noise-corrupted Image



< Original
image (left)

Noisy image >
(right)
(Gaussian
noise)



Noise Reduction Using Linear Filtering



< Original
grayscale
image (left)

Noisy image >
(right) (Salt
and pepper
noise, pretty
common for
old b&w
cameras)



Noise Reduction Using Linear Filtering



Gaussian filter
used to
denoise both
noisy images
(25 x 25) $\sigma^2=3$

Output is
blurred, s&p
not properly
removed



Noise Reduction Using Nonlinear Filtering



Median filter is
used to
denoise the
salt and
pepper noise

(can you tell
the original
from the
filtered one?)




Unsharp Masking & High Boost Filtering

This unsharp masking process consists of the following steps:

1. Blur the original image
 - a. $f'(x, y) = \text{Blur}[f(x, y)]$
2. Subtract the blurred image from the original (the resulting difference is called the mask)
 - a. $g_{\text{mask}}(x, y) = f(x, y) - f'(x, y)$
3. Add the mask to the original
 - a. $g(x, y) = f(x, y) + k * g_{\text{mask}}(x, y)$

$k = 1$ is *unsharp masking*, $k > 1$ is *high boost filtering*



Unsharp Masking

Original image



Unsharp masking



High Boost Filtering

Original image



High boost filtering



Geometric Transformations

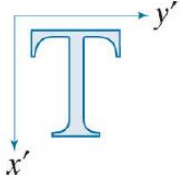
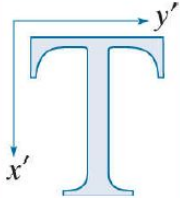
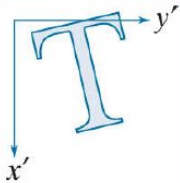
- We use geometric transformations modify the spatial arrangement of pixels in an image
- These transformations are called rubber-sheet transformations because they may be viewed as analogous to “printing” an image on a rubber sheet, then stretching or shrinking the sheet according to a predefined set of rules

Geometric Transformations

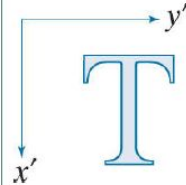
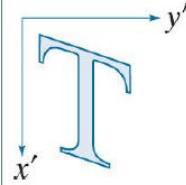
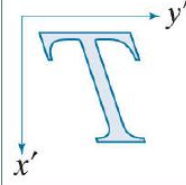
- (x, y) are pixel coordinates in the original image
- (x', y') are the corresponding pixel coordinates of the transformed image
- Our interest is in so-called affine transformations, which include scaling, translation, rotation, and shearing
- The key characteristic of an affine transformation in 2-D is that it preserves points, straight lines, and planes
-

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \mathbf{A} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Geometric Transformations

Transformation Name	Affine Matrix, A	Coordinate Equations	Example
Identity	$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$	$\begin{aligned} x' &= x \\ y' &= y \end{aligned}$	
Scaling/Reflection (For reflection, set one scaling factor to -1 and the other to 0)	$\begin{bmatrix} c_x & 0 & 0 \\ 0 & c_y & 0 \\ 0 & 0 & 1 \end{bmatrix}$	$\begin{aligned} x' &= c_x x \\ y' &= c_y y \end{aligned}$	
Rotation (about the origin)	$\begin{bmatrix} \cos\theta & -\sin\theta & 0 \\ \sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$	$\begin{aligned} x' &= x\cos\theta - y\sin\theta \\ y' &= x\sin\theta + y\cos\theta \end{aligned}$	

Geometric Transformations

Transformation Name	Affine Matrix, A	Coordinate Equations	Example
Translation	$\begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix}$	$x' = x + t_x$ $y' = y + t_y$	
Shear (vertical)	$\begin{bmatrix} 1 & s_v & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$	$x' = x + s_v y$ $y' = y$	
Shear (horizontal)	$\begin{bmatrix} 1 & 0 & 0 \\ s_h & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$	$x' = x$ $y' = s_h x + y$	

Tutorial

Basic python implementation of some of the algorithms taught in class can be found in the following google colab notebook:

<https://colab.research.google.com/drive/1JF3ww3l-bBzoC0iy9Ado7eQLcxC1hVFU?usp=sharing>

Thank you!